

1. The “Installing the control plane” section (<https://www.ibm.com/docs/en/cloud-paks/cp-data/3.5.0?topic=installing-control-plane>) should include information on the known issue and workaround for inode-limit for PVCs:

There is a problem with the inode-limit for the PVCs created during the install in which the set inode-limit was too low for the `user-home-pvc` (defaulted in our environment to 3072 which was dynamically calculated using `volume size / block size of the filesystem`), but at least 4.4K is required for CP4D Install. A workaround is to adjust the inode-limit after the `user-home-pv` is created and bound during the installation.



There is a problem with the inode-limit for the PVCs created during the install in which the set inode-limit was too low for the `user-home-pvc` (defaulted in our environment to 3072 which was dynamically calculated using `volume size / block size of the filesystem`), but at least 4.4K is required for CP4D Install. A workaround is to adjust the inode-limit after the `user-home-pv` is created and bound during the installation.

<https://wiki.cpst-lab.no-users.ibm.com/en/general/CP4D>

2. On the “Installing from cpd-cli” section (<https://www.ibm.com/docs/en/cloud-paks/cp-data/3.5.0?topic=plane-installing-from-cpd-cli>) it would also be helpful to add a step after step #4 about how to verify creation of PVC. For example, the following could be added before the “Verifying that the installation completed successfully section”:
<https://wiki.cpst-lab.no-users.ibm.com/en/general/CP4D>

► Watch for the creation of the `user-home-pvc` with the command

```
1 | oc get pvc -n zen -w
```



Wait for the creation of the PVC and for STATUS of Bound.

```
1 | [root@arcx3650fxxnh ~]# oc get pvc user-home-pvc -n zen
2 | NAME          STATUS    VOLUME                                     CAPACITY   ACCESS MODES   STORAGECLASS   AGE
3 | user-home-pvc  Bound    pvc-90346e43-cbf5-467c-b83e-c93a4832eb5e  10Gi       RWX            ibm-spectrum-scale-sc  16h
```



Note the pvc volume for `user-home-pvc` (in this case, `pvc-90346e43-cbf5-467c-b83e-c93a4832eb5e`).

► On the remote storage cluster, adjust the inode-limit for the PVC

```
1 | mmchfileset fs1 pvc-90346e43-cbf5-467c-b83e-c93a4832eb5e --inode-limit '6144:6144'
```

Sample output

```
1 | [root@stg-node0 ~]# mmchfileset fs1 pvc-90346e43-cbf5-467c-b83e-c93a4832eb5e --inode-limit '6144:6144'
2 | Set maxInodes for inode space 8 to 6144
3 | Fileset pvc-90346e43-cbf5-467c-b83e-c93a4832eb5e changed.
4 | [root@stg-node0 ~]# mmlsfileset fs1 pvc-90346e43-cbf5-467c-b83e-c93a4832eb5e -l
5 | Collecting fileset usage information ...
6 | Filesets in file system 'fs1':
7 | Name          Status    Path                                     InodeSpace   MaxInodes   AllocInodes   UsedInodes
8 | pvc-90346e43-cbf5-467c-b83e-c93a4832eb5e Linked /ibm/fs1/pvc-90346e43-cbf5-467c-b83e-c93a4832eb5e 8      6144      6144      20
```

- Under the “Verifying that the installation completed successfully” section, the `./cpd-cli status` command does not include the `--repo` flag (<https://www.ibm.com/docs/en/cloud-paks/cp-data/3.5.0?topic=plane-installing-from-cpd-cli>)

Verifying that the installation completed successfully

From your installation node:

- Run the following command:

```
./cpd-cli status \  
--assembly lite \  
--namespace Project
```

Replace *Project* with the value you used in when you installed control plane.

- If the installation completed successfully, the status of the assembly and the modules in the assembly is Ready.
- If the installation failed, contact IBM Support for assistance.

Here is an example of the `./cpd-cli status` command from our CPST team wiki:

Installation status

► View installation status

```
1 | ./cpd-cli status --repo ./repo.yaml --namespace zen --assembly lite
```

- Under the “Installing the Cloud Pak for Data control plane from cpd-cli”, it would be useful to include a section about the 4MiB file system workaround. <https://wiki.cpst-lab.no-users.ibm.com/en/general/CP4D>

This workaround is only necessary if the storage cluster file system is set to 4 MiB block size (Scale GUI defaults to 4 MiB).

Watch for the creation of the `user-home-pvc` with the command:

```
oc get pvc -n zen -w
```

Wait for the creation of the PVC and for STATUS of Bound.

```
[root@arcx3650fxnh ~]# oc get pvc user-home-pvc -n zen  
NAME          STATUS  VOLUME                                     CAPACITY  ACCESS MODES  
STORAGECLASS  AGE  
user-home-pvc Bound   pvc-90346e43-cbf5-467c-b83e-c93a4832eb5e 10Gi      RWX  
ibm-spectrum-scale-sc 16h
```

Note the pvc volume for `user-home-pvc` (in this case, **pvc-90346e43-cbf5-467c-b83e-c93a4832eb5e**).

On the remote storage cluster, adjust the inode-limit for the PVC:

```
mmchfileset fs1 pvc-90346e43-cbf5-467c-b83e-c93a4832eb5e --inode-limit '6144:6144'
```

5. In the “Creating a custom route to the platform section” (<https://www.ibm.com/docs/en/cloud-paks/cp-data/3.5.0?topic=tasks-creating-custom-route-platform>), it would be useful to also list the command to verify available routes after step #3.

The command to verify available routes is: `oc get route`

Verify available routes

```
1 [root@arcx3650fxxnh ~]# oc get route
2 NAME          HOST/PORT          PATH    SERVICES    PORT    TERMINATION
3 cp4d-cpst      cp4d-cpst-zen.apps.cpst-ocp-cluster-a.cpst-lab.no-users.ibm.com  ibm-nginx-svc  ibm-nginx-https-port  passthrough
4 zen-cpd        zen-cpd-zen.apps.cpst-ocp-cluster-a.cpst-lab.no-users.ibm.com    ibm-nginx-svc  ibm-nginx-https-port  passthrough/Redirect
```

<https://wiki.cpst-lab.no-users.ibm.com/en/general/CP4D>

6. After step 5 in the “Preparing clusters connected to the internet section” <https://www.ibm.com/docs/en/cloud-paks/cp-data/3.5.0?topic=db2-setting-up-cluster> it would be useful to also include information about how to grant cpd-admin-role to the project administrator user:

For example:

Run the following command to grant cpd-admin-role to the project administration user:

```
oc adm policy add-role-to-user cpd-admin-role ocpadmin --role-namespace=zen -n zen
```

<https://wiki.cpst-lab.no-users.ibm.com/en/general/db2oltp>

7. Under the ‘Setting up the cluster for Db2 Warehouse’, after step 5 in the “Preparing clusters connected to the internet section” <https://www.ibm.com/docs/en/cloud-paks/cp-data/3.5.0?topic=warehouse-setting-up-cluster-db2> it would be useful to also include information about how to grant cpd-admin-role to the project administrator user:

For example:

Run the following command to grant cpd-admin-role to the project administration user:

```
oc adm policy add-role-to-user cpd-admin-role ocpadmin --role-namespace=zen -n zen
```

<https://wiki.cpst-lab.no-users.ibm.com/en/general/db2wh>

8. Under the ‘Setting up the cluster for DataStage Enterprise’, after step 4 in the “Preparing clusters connected to the internet section” (<https://www.ibm.com/docs/en/cloud-paks/cp-data/3.5.0?topic=enterprise-setting-up-cluster-datastage>) it would be useful to also include information about how to grant cpd-admin-role to the project administrator user:

For example:

Run the following command to grant cpd-admin-role to the project administration user:

```
oc adm policy add-role-to-user cpd-admin-role ocpadmin --role-namespace=zen-automated -n zen-automated
```

<https://wiki.cpst-lab.no-users.ibm.com/en/general/DataStage>

9. In the “Installing The DataStage Enterprise service” section before “What to do next” it would be beneficial to include the following Post Installation tasks

Data Stage post installation tasks

Edit the statefulset

```
oc edit sts/is-en-conductor
```

Insert the following entry under VolumeMounts

```
VolumeMounts:
  mountPath: /home/dsadm
  name: engine-dedicated-volume
  subPath: is-en-conductor-0/EngineClients/db2_client/dsadm
```

Upon saving, the pod will restart

There is an extra step here, but only if the pod has been restarted since installing (or if the files put in /home/dsadm by default are otherwise deleted)
Files that should be in /home/dsadm

```

1 [root@arcx3650fxxnh ~]# oc exec -n zen-automated is-en-conductor-0 -- ls -al /home/dsadm
2 total 204
3 drwxrwx--x. 5 root root 4096 Jun 3 17:19 .
4 drwxr-xr-x. 1 root root 51 May 5 04:06 ..
5 -rw----- 1 dsadm dstage 55 Jun 3 16:21 .bash_history
6 -rwxr-xr-x. 1 dsadm dstage 18 Aug 21 2019 .bash_logout
7 -rwxr-xr-x. 1 dsadm dstage 193 Aug 21 2019 .bash_profile
8 -rwxr-xr-x. 1 dsadm dstage 344 May 5 04:06 .bashrc
9 drwxr-xr-x. 2 dsadm dstage 4096 May 5 03:47 ds_logs
10 -rw-r--r-- 1 dsadm dstage 0 Jun 3 17:19 .extractComplete
11 drwxr-xr-x. 2 dsadm dstage 4096 Jun 3 17:22 imam_logs
12 drwxrw---- 3 dsadm dstage 4096 Jun 3 00:01 .pki

```

If any are missing, delete `.extractComplete`, then restart the pod:

```
oc exec -n zen-automated is-en-conductor-0 -- rm -f /home/dsadm/.extractComplete
```

10. In the “Uninstalling DataStage Enterprise” Section under the “What to do next” topic it would be beneficial to include the following:

Run the following command if pvc or pv get hung “Terminating” status:

```
oc patch pvc {PVC_NAME} -p '{"metadata":{"finalizers":null}}'
```

```
oc patch pv {PVC_NAME} -p '{"metadata":{"finalizers":null}}'
```

11. In the “Uninstalling DataStage Enterprise” and “Uninstalling DataStage Enterprise Plus” sections, the commands to uninstall DataStage should have the `--profile` flag when they are run

```
./cpd-cli uninstall \
--assembly ds \
--namespace zen-automated \
--profile cpst-test-profile
--dry-run
```

```
./cpd-cli uninstall \
--assembly ds \
--namespace zen-automated \
--profile cpst-test-profile
```

12. Also under the “Uninstalling DataStage Enterprise” section the following should be added:

Also remove the iis component (assuming its not used by other services) for a true uninstall

```
./cpd-cli uninstall \
--assembly iis \
--namespace zen-automated \
--profile cpst-test-profile
```

13. Under the “Uninstalling DataStage Enterprise” section the commands in the “Delete all PersistentVolumeClaim (PVC) and PersistentVolume (PV) resources:”, “Delete all job resources:” and “Delete all cronjob resources:” sections should be replaced with the following commands (which need to be run in this specific order):

```

oc get deployment --sort-by='{.metadata.creationTimestamp}' | grep ds-'
oc delete deployment <deployment_name>
oc get jobs --sort-by='{.metadata.creationTimestamp}' | grep ds-'
oc delete job <job_name>
oc get StatefulSet --sort-by='{.metadata.creationTimestamp}' | grep ds-'
oc delete StatefulSet <StatefulSet_name>
oc get ReplicaSet --sort-by='{.metadata.creationTimestamp}' | grep ds-'
oc delete ReplicaSet <ReplicaSet_name>
oc get services --sort-by='{.metadata.creationTimestamp}' | grep ds-'
oc delete service <service_name>
oc get pods --sort-by='{.metadata.creationTimestamp}' | grep ds-'
oc delete pod <pod_name>
oc get pvc --sort-by='{.metadata.creationTimestamp}' | grep ds-'
oc delete pvc <pvc_name>
oc get pv --sort-by='{.metadata.creationTimestamp}' | grep ds-'
oc delete pv <pv_name>

```

14. In the “Installing Watson Machine Learning” section, the command to install the wml service could contain the --target-registry-username and --target-registry-password flags depending on how the cluster has been configured by the customer

```

./cpd-cli install \
--repo ./repo.yaml \
--assembly wml \
--arch x86_64 \
--namespace zen \
--storageclass ibm-spectrum-scale-sc \
--transfer-image-to-registry.cpst-lab.no-users.ibm.com \
--cluster-pull-prefix registry.cpst-lab.no-users.ibm.com \
--ask-push-registry-credentials \
--latest-dependency \
--dry-run

```

```

Username: oc whoami
Password: oc whoami -t

```

15. In the “Uninstalling Watson Machine Learning” section, the command to uninstall wml should have the --profile flag when it is run

```

./cpd-cli uninstall \
--assembly wml \
--namespace zen \
--profile cpst-test-profile

```

16. In the “Installing Watson Studio” section, the command to install the wsl service could contain the --target-registry-username and --target-registry-password flags depending on how the cluster has been configured by the customer

```

./cpd-cli install \
--repo ./repo.yaml \
--assembly wsl \
--arch x86_64 \
--namespace zen \
--storageclass ibm-spectrum-scale-sc \
--transfer-image-to-registry.cpst-lab.no-users.ibm.com \
--cluster-pull-prefix registry.cpst-lab.no-users.ibm.com \
--ask-push-registry-credentials \

```

```
--latest-dependency \  
--dry-run
```

```
Username: oc whoami  
Password: oc whoami -t
```

17. In the “Uninstalling Watson Studio” section, the command to uninstall wsl should have the `--profile` flag when it is run

```
./cpd-cli uninstall \  
--assembly wsl \  
--namespace zen \  
--profile cpst-test-profile
```

18. In the section “Preparing the cluster for the Watson Knowledge Catalog service” prior to the “Next Step” section, it would be beneficial to add the following:

Run the following command to grant `cpd-admin-role` to the project administration user:

```
oc adm policy add-role-to-user cpd-admin-role ocpadmin --role-namespace=zen-  
automated -n zen-automated
```

19. In the “Uninstalling Watson Knowledge Catalog” section, the command to uninstall wkc should have the `--profile` flag when it is run

```
./cpd-cli uninstall \  
--assembly wkc \  
--namespace zen \  
--profile cpst-test-profile
```

20. After the uninstall command is run, the following commands should be added in order to make sure everything is removed. The commands should be run in the following order:

```
oc get deployment --sort-by='{.metadata.creationTimestamp}'  
oc delete deployment <deployment_name>
```

```
oc get jobs --sort-by='{.metadata.creationTimestamp}'  
oc delete job <job_name>
```

```
oc get StatefulSet --sort-by='{.metadata.creationTimestamp}'  
oc delete StatefulSet <StatefulSet_name>
```

```
oc get ReplicaSet --sort-by='{.metadata.creationTimestamp}'  
oc delete ReplicaSet <ReplicaSet_name>
```

```
oc get services --sort-by='{.metadata.creationTimestamp}'  
oc delete service <service_name>
```

```
oc get pods --sort-by='{.metadata.creationTimestamp}'  
oc delete pod <pod_name>
```

```
oc get pvc --sort-by='{.metadata.creationTimestamp}'  
oc delete pvc <pvc_name>
```

```
oc get pv --sort-by='{.metadata.creationTimestamp}'
oc delete pv <pv_name>

oc get configmap --sort-by='{.metadata.creationTimestamp}'
oc delete configmap <configmap_name>
```

21. After running the commands above commands to remove all instances belonging to Watson Knowledge Catalog, the following commands should be run in case pvc or pv get hung

```
oc patch pvc {PVC_NAME} -p '{"metadata":{"finalizers":null}}'
oc patch pv {PVC_NAME} -p '{"metadata":{"finalizers":null}}'
```

22. In the “Installing Data Virtualization” section, it would be beneficial to include information regarding checking for patches after running the command under “Verifying that the installation completed successfully”

Checking for available patches

```
./cpd-cli status \
--repo ./repo.yaml \
--namespace zen \
--assembly dv \
--patches \
--available-updates
```

23. In the “Installing DB2 Data Management Console” section, it would be beneficial to include information regarding checking for patches after running the command under “Verifying that the installation completed successfully”

Checking for available patches

```
./cpd-cli status \
--repo ./repo.yaml \
--namespace zen \
--assembly dmc \
--patches \
--available-updates
```

24. The “Uninstalling DB2 Data Management Console” section is missing the procedures to uninstall dmc

Run the following command to preview what will be removed for dmc

```
./cpd-cli uninstall \
--assembly dmc \
--namespace zen \
--profile cpst-test-profile
--dry-run
```

25. Under the ‘Setting up the cluster for Analytics Engine Powered by Apache Spark’, before the ‘What to do next’ section, it would be useful to also include information about how to grant cpd-admin-role to the project administrator user:

For example:

Run the following command to grant cpd-admin-role to the project administration user:

```
oc adm policy add-role-to-user cpd-admin-role ocpadmin --role-namespace=zen -n zen
```

Future Documentation Enhancements:

26. Once CP4D officially supports CNSA/CSI storage classes, the section below could be updated to mention how to set up storage class for Spectrum Scale.

<https://www.ibm.com/docs/en/cloud-paks/cp-data/3.5.0?topic=planning-storage-considerations>

The screenshot shows the IBM Cloud Pak for Data 3.5.0 documentation page. The sidebar on the left includes the IBM logo, 'Documentation', a search bar, and a table of contents. The 'Storage considerations' section is highlighted in the table of contents. The main content area displays a table with two rows: 'Storage classes' and 'Data replication for high availability'. The 'Storage classes' row lists required storage classes, user-defined NFS storage classes, and a table of storage classes (ibmc-file-gold-gid, ocs-storagecluster-cephfs, ocs-storagecluster-ceph-rbd). The 'Data replication for high availability' row lists supported replication options, including supported by default, supported by default with 3 replicas, and supported but not enabled by default.

Storage classes	The required storage classes are automatically created when you install OpenShift Container Storage. Cloud Pak for Data uses the following storage classes:	NFS storage classes are user-defined. Use a storage class with ReadWriteMany (RWX) access.	The required storage classes are listed in Creating Portworx storage classes . You can run the provided script to create the storage classes.	ibmc-file-gold-gid
	<ul style="list-style-type: none">ocs-storagecluster-cephfsocs-storagecluster-ceph-rbd			
Data replication for high availability	Supported By default, all services use multiple replicas for high availability. OpenShift Container Storage maintains each replica in a distinct availability zone.	Replication support depends on your NFS server.	Supported By default, most services use a storage class that supports 3 replicas. For details about the replicas for each storage class, see Creating Portworx storage classes .	Supported, but not enabled by default. You can enable replication from the IBM Cloud console. For details, see Replicating data .

The steps below, could be integrated into the CSM Playbook.

<https://wiki.cpst-lab.no-users.ibm.com/en/general/CP4D>

Search...







TABLE OF CONTENTS

> Pre-installation Tasks

> Setting up a storage class for CP4D

> Obtaining the installation files.

> Set up the requirements for the cpd-cli command

> Setting up the cluster for the control plane

> Installing the control plane

> Installation status

> Post-installation tasks


> Creating a custom route to the platform

> Setting up the web client

> Uninstalling Cloud Pak for Data


TAGS

public




TALK

View Discussion



LAST EDITED BY

Jayson E Tsingine



Pre-installation Tasks

Setting up a storage class for CP4D

For the CPST environment, we will be using *dynamic provisioning* of persistent volumes (PVs) based on *independent filesets* in IBM Spectrum Scale.

> Create `ibm-spectrum-scale-sc.yaml` file with the following content:

```
1 | apiVersion: storage.k8s.io/v1
2 | kind: StorageClass
3 | metadata:
4 |   name: ibm-spectrum-scale-sc
5 | provisioner: spectrumscale.csi.ibm.com
6 | parameters:
7 |   volBackendFs: "fs1"
8 |   clusterId: "10724741469548137692"
9 | reclaimPolicy: Delete
```

> Create the storage class

```
1 | oc create -f ibm-spectrum-scale-sc.yaml
```

> Set the storage class as the default storage class

```
1 | oc patch storageclass ibm-spectrum-scale-sc -p '{"metadata": {"annotations":{"storageclass.kubernetes.io/is-default-class":"true"}}}'
```

> View storage class

```
1 | [root@arcx3650fxxnh ~]# oc get storageclass ibm-spectrum-scale-sc
2 | NAME                                PROVISIONER      RECLAIMPOLICY   VOLUMEBINDINGMODE   ALLOWVOLUMEEXPANSION   AGE
3 | ibm-spectrum-scale-sc (default)    spectrumscale.csi.ibm.com   Delete          Immediate            false                  8d
```

